

Inference Energy by Model Variant

CodeCarbon modeled estimates; normalized from a fixed 100-example inference workload.

Estimated energy (Wh per 1,000 predictions, log scale)

10²

10¹

10⁰

0.141

26.089

29.747

34.387

74.510

Distilled DistilBERT

Qwen FP16

Qwen Pruned 30%

Qwen INT4 NF4

Qwen INT8

